

```
In [2]: !pip3 install xgboost
```

Requirement already satisfied: xgboost in /usr/local/lib/python3.7/dist-packages (0.90)

Requirement already satisfied: scipy in /usr/local/lib/python3.7/dist-packages (from xgboost) (1.4.1)

Requirement already satisfied: numpy in /usr/local/lib/python3.7/dist-packages (from xgboost) (1.19.5)

```
In [180]: #Importing Libraries
# please do go through this python notebook:
import warnings
warnings.filterwarnings("ignore")

import csv
import pandas as pd#pandas to create small dataframes
import datetime #Convert to unix time
import time #Convert to unix time
# if numpy is not installed already : pip3 install numpy
import numpy as np#Do arithmetic operations on arrays
# matplotlib: used to plot graphs
import matplotlib
import matplotlib.pyplot as plt
import seaborn as sns#Plots
from matplotlib import rcParams#Size of plots
from sklearn.cluster import MiniBatchKMeans, KMeans#Clustering
import math
import pickle
import os
# to install xgboost: pip3 install xgboost
import xgboost as xgb

import warnings
import networkx as nx
import pdb
import pickle
from pandas import HDFStore, DataFrame
from pandas import read_hdf
from scipy.sparse.linalg import svds, eigs
import gc
from tqdm import tqdm
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import f1_score

from sklearn.model_selection import GridSearchCV

from plotly.offline import iplot
import plotly.graph_objs as go
```

```
In [5]: !gdown --id "1fDJptlCFEWNV5UNGpc4geTykgFI3PDCV"
!gdown --id "1pIO_nOg9XU0WUD10brRvrgyUXbY5gqMs"
!gdown --id "1lcxzVZ0-MkPmoH3lS35Q8rRfrecKSxb1"
!gdown --id "1_KN7S8zfHdRkRjRYOEtBxBVq8JrGxPXD"
```

Downloading...

From: <https://drive.google.com/uc?id=1fDJptlCFEWNV5UNGpc4geTykgFI3PDCV>

To: /content/storage\_sample\_stage4.h5

103MB [00:01, 65.6MB/s]

```
In [11]: from pandas import read_hdf
df_train = read_hdf('storage_sample_stage4.h5', 'train_df', mode='r')
df_test = read_hdf('storage_sample_stage4.h5', 'test_df', mode='r')
```

## Featurization

### Add missing Feature : num\_followers\_d

```
In [103]: train_graph=nx.read_edgelist('train_after_eda.csv',delimiter=',',create_using=
nx.DiGraph(),nodetype=int)
```

```
In [104]: train_d_followers = [len(set(train_graph.predecessors(dest_node))) for dest_no
de in df_train["destination_node"].values]
test_d_followers = [len(set(train_graph.predecessors(dest_node))) for dest_nod
e in df_test["destination_node"].values]
```

```
In [106]: df_train["num_followers_d"] = train_d_followers
df_test["num_followers_d"] = test_d_followers
```

### Add Feature : Preferential Attachment

```
In [120]: df_train["pref_attach_followees"] = df_train["num_followees_s"]*df_train["num_
followees_d"]
df_train["pref_attach_followers"] = df_train["num_followers_s"]*df_train["num_
followers_d"]

df_test["pref_attach_followees"] = df_test["num_followees_s"]*df_test["num_fol
lowees_d"]
df_test["pref_attach_followers"] = df_test["num_followers_s"]*df_test["num_fol
lowers_d"]
```

### Add Feature : svd\_dot

```
In [109]: def compute_svd_dot(df_row):
    svd_dot_u = (df_row["svd_u_s_1"]*df_row["svd_u_d_1"]) + \
        (df_row["svd_u_s_2"]*df_row["svd_u_d_2"]) + \
        (df_row["svd_u_s_3"]*df_row["svd_u_d_3"]) + \
        (df_row["svd_u_s_4"]*df_row["svd_u_d_4"]) + \
        (df_row["svd_u_s_5"]*df_row["svd_u_d_5"]) + \
        (df_row["svd_u_s_6"]*df_row["svd_u_d_6"])

    svd_dot_v = (df_row["svd_v_s_1"]*df_row["svd_v_d_1"]) + \
        (df_row["svd_v_s_2"]*df_row["svd_v_d_2"]) + \
        (df_row["svd_v_s_3"]*df_row["svd_v_d_3"]) + \
        (df_row["svd_v_s_4"]*df_row["svd_v_d_4"]) + \
        (df_row["svd_v_s_5"]*df_row["svd_v_d_5"]) + \
        (df_row["svd_v_s_6"]*df_row["svd_v_d_6"])

    return svd_dot_u, svd_dot_v
```

```
In [126]: for index, row in tqdm(df_train.iterrows()):
    df_train.loc[index,"svd_dot_u"] ,df_train.loc[index,"svd_dot_v"]= compute_svd_dot(row)

    for index, row in tqdm(df_test.iterrows()):
        df_test.loc[index,"svd_dot_u"] ,df_test.loc[index,"svd_dot_v"] = compute_svd_dot(row)
```


100002it [03:00, 554.21it/s]

50002it [01:03, 789.98it/s]

```
In [127]: df_test.head(1)
```

Out[127]:

	source_node	destination_node	indicator_link	jaccard_followers	jaccard_followees	cosine_fol
0	848424	784690	1	0	0.0	0.0



```
In [128]: # Save final train and test df
df_train.to_csv("train_data.csv",index=False)
df_test.to_csv("test_data.csv",index=False)
```

## XGBOOST Algorithm

```
In [153]: X_train = pd.read_csv("train_data.csv")
X_test = pd.read_csv("test_data.csv")
y_train = X_train.indicator_link
y_test = X_test.indicator_link

X_train.drop(['source_node', 'destination_node', 'indicator_link'],axis=1,inplace=True)
X_test.drop(['source_node', 'destination_node', 'indicator_link'],axis=1,inplace=True)

X_train = X_train.values
X_test = X_test.values
y_train = y_train.values
y_test = y_test.values
```

## Hyperparameter Tuning

```
In [186]: clf_xgb = xgb.XGBClassifier(objective = 'binary:logistic')

param_dist = {
    'n_estimators': [1,5, 10],
    'max_depth': [3, 5, 8],
}

hyp_clf = GridSearchCV(clf_xgb, param_dist,cv=5, scoring = 'f1', return_train_score = True)
```

```
In [187]: hyp_clf.fit(X_train,y_train)
```

```
Out[187]: GridSearchCV(cv=5, error_score=nan,
                      estimator=XGBClassifier(base_score=0.5, booster='gbtree',
                                              colsample_bylevel=1, colsample_bynode=1,
                                              colsample_bytree=1, gamma=0,
                                              learning_rate=0.1, max_delta_step=0,
                                              max_depth=3, min_child_weight=1,
                                              missing=None, n_estimators=100, n_jobs=
1,
                                              nthread=None, objective='binary:logisti
c',
                                              random_state=0, reg_alpha=0, reg_lambda=
1,
                                              scale_pos_weight=1, seed=None, silent=No
ne,
                                              subsample=1, verbosity=1),
                      iid='deprecated', n_jobs=None,
                      param_grid={'max_depth': [3, 5, 8], 'n_estimators': [1, 5, 10]},
                      pre_dispatch='2*n_jobs', refit=True, return_train_score=True,
                      scoring='f1', verbose=0)
```

```
In [188]: best_depth = hyp_clf.best_estimator_.max_depth
best_n = hyp_clf.best_estimator_.n_estimators

depth_vals = hyp_clf.cv_results_["param_max_depth"]
n_estimators_vals = hyp_clf.cv_results_["param_n_estimators"]
auc_vals_cv = hyp_clf.cv_results_["mean_test_score"]
auc_vals_train = hyp_clf.cv_results_["mean_train_score"]
```

```
In [191]: print(best_depth, best_n)
```

8 10

```
In [190]: #3d plot : max_depth vs n_estimators vs f1_score
trace1 = go.Scatter3d(x=depth_vals,y=n_estimators_vals,z=auc_vals_train, name = 'train')
trace2 = go.Scatter3d(x=depth_vals,y=n_estimators_vals,z=auc_vals_cv, name = 'Cross validation')
data = [trace1, trace2]
layout = go.Layout(scene = dict(xaxis = dict(title='max_depth'),
                                yaxis = dict(title='n_estimators'),
                                zaxis = dict(title='f1_Score')),)
fig = go.Figure(data=data, layout=layout)
iplot(fig, filename='3d-scatter-colorscale')
```

## Model Building

```
In [198]: best_model = xgb.XGBClassifier(objective = 'binary:logistic',max_depth=best_de  
pth,n_estimators = best_n)  
best_model.fit(X_train, y_train,eval_set=[(X_test, y_test)],eval_metric='auc')
```

```
[0]    validation_0-auc:0.954186  
[1]    validation_0-auc:0.954272  
[2]    validation_0-auc:0.953754  
[3]    validation_0-auc:0.953282  
[4]    validation_0-auc:0.953875  
[5]    validation_0-auc:0.95379  
[6]    validation_0-auc:0.95353  
[7]    validation_0-auc:0.954014  
[8]    validation_0-auc:0.954407  
[9]    validation_0-auc:0.954795
```

```
Out[198]: XGBClassifier(base_score=0.5, booster='gbtree', colsample_bylevel=1,  
    colsample_bynode=1, colsample_bytree=1, gamma=0,  
    learning_rate=0.1, max_delta_step=0, max_depth=8,  
    min_child_weight=1, missing=None, n_estimators=10, n_jobs=1,  
    nthread=None, objective='binary:logistic', random_state=0,  
    reg_alpha=0, reg_lambda=1, scale_pos_weight=1, seed=None,  
    silent=None, subsample=1, verbosity=1)
```

```
In [200]: y_train_pred = best_model.predict(X_train)  
y_test_pred = best_model.predict(X_test)
```

## Evaluating Results

```
In [201]: from sklearn.metrics import f1_score  
print('Train f1 score',f1_score(y_train,y_train_pred))  
print('Test f1 score',f1_score(y_test,y_test_pred))
```

```
Train f1 score 0.9718951457100442  
Test f1 score 0.9333727660653175
```

```
In [202]: from sklearn.metrics import confusion_matrix
def plot_confusion_matrix(test_y, predict_y):
    C = confusion_matrix(test_y, predict_y)

    A = ((C.T)/(C.sum(axis=1))).T

    B = (C/C.sum(axis=0))
    plt.figure(figsize=(20,4))

    labels = [0,1]
    # representing A in heatmap format
    cmap=sns.light_palette("blue")
    plt.subplot(1, 3, 1)
    sns.heatmap(C, annot=True, cmap=cmap, fmt=".3f", xticklabels=labels, ytick
labels=labels)
    plt.xlabel('Predicted Class')
    plt.ylabel('Original Class')
    plt.title("Confusion matrix")


    plt.subplot(1, 3, 2)
    sns.heatmap(B, annot=True, cmap=cmap, fmt=".3f", xticklabels=labels, ytick
labels=labels)
    plt.xlabel('Predicted Class')
    plt.ylabel('Original Class')
    plt.title("Precision matrix")

    plt.subplot(1, 3, 3)
    # representing B in heatmap format
    sns.heatmap(A, annot=True, cmap=cmap, fmt=".3f", xticklabels=labels, ytick
labels=labels)
    plt.xlabel('Predicted Class')
    plt.ylabel('Original Class')
    plt.title("Recall matrix")


    plt.show()
```

```
In [203]: print('Train confusion_matrix')
plot_confusion_matrix(y_train,y_train_pred)
print('Test confusion_matrix')
plot_confusion_matrix(y_test,y_test_pred)
```


Train confusion\_matrix



Test confusion\_matrix




```
In [204]: from sklearn.metrics import roc_curve, auc
fpr,tpr,ths = roc_curve(y_test,y_test_pred)
auc_sc = auc(fpr, tpr)
plt.plot(fpr, tpr, color='navy',label='ROC curve (area = %0.2f)' % auc_sc)
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic with test data')
plt.legend()
plt.show()
```





```
In [207]: features = df_train.columns
importances = best_model.feature_importances_
indices = (np.argsort(importances))[-25:]
plt.figure(figsize=(10,12))
plt.title('Feature Importances')
plt.barh(range(len(indices)), importances[indices], color='r', align='center')
plt.yticks(range(len(indices)), [features[i] for i in indices])
plt.xlabel('Relative Importance')
plt.show()
```



In [206]:

## Conclusion

## STEPS Followed in the Case study

- Dataset Generation
  - Reading the dataset which is in the form  $\langle \text{src\_id}, \text{dest\_id} \rangle$ . Let the indicator for the read points be = 1 (indicates src is connected to dest)
  - Using the read dataset, create a Graph using the Nodes provided and each  $\langle \text{src\_id}, \text{dest\_id} \rangle$  indicating an edge in the graph
  - Generate random  $\langle \text{src}, \text{dest} \rangle$  pairs for which  $\text{shortestPath}(\text{src}, \text{dest}) > 2$  and combine it with originally read dataset. The value of indicator is 0 (indicating src and dest are not connected)
- Once dataset is generated, split it into Train Data (75%) and Test Data (25%)

## Featurization Part :

The following features are calculated and added as features -Followers and Followees of src and dest

- Similarity Measures : Jaccards Coefficient, Otsuka-Ochia Coefficient
- Page Rank
- Whether Src and dest belong to same Weakly Connected Component
- Adar Index
- Katz Centrality
- HITS
- Use singular value decomposition to get U and V.T each of size (1,6) for src and dest
- Weight Features
- Preferential Attachment
- SVD Dot features

## Model Building

- The algorithm used here was XGBoost algorithm.
- Hyperparameter tuning was done using Grid Search CV and the best parameters were found to be
- $\text{max\_depth}=8, \text{n\_estimators} = 10$

## Final Results :

- Train f1 score 0.9718951457100442
- Test f1 score 0.9333727660653175

In [ ]: